



# EE 046202 - Technion - Unsupervised Learning & Data Analysis

- Formerly 046193

Tal Daniel

## Tutorial 06 - Dimensionality Reduction - Kernels & KPCA



### Agenda

- Recap: PCA Algorithm
- Kernels
  - The Kernel Trick
- NLPKA- Nonlinear PCA
  - Kernel PCA

```
In [1]: # imports for the tutorial
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from sklearn.decomposition import PCA, KernelPCA
from sklearn.datasets import make_circles, make_moons
%matplotlib notebook
```



### PCA Algorithm

1. **Normalize/Standardize** (if we use features of different scales, we may get misleading components) and center the data. Given data  $X \in \mathcal{R}^{m \times N}$ , where  $m$  is the number of features and  $N$  is the number of samples, normalization:

$$\tilde{X} = X - \bar{X}$$

Standardization:

$$\tilde{X} = \frac{X - \bar{X}}{\bar{\sigma}_x}$$

Where  $\bar{\sigma}_x$  is the empirical standard deviation (the square root of the empirical variance).

2. Calculate the covariance matrix  $\tilde{X}$  of data points:

$$P = \tilde{X}\tilde{X}^T \in \mathcal{R}^{m \times m}$$

3. Calculate eigenvectors and corresponding eigenvalues.
4. Sort the eigenvectors according to their eigenvalues in decreasing order.
5. Choose first  $k$  largest eigenvectors and that will be the new  $k$  dimensions.
6. Transform the original  $d$  dimensional data points into  $k$  dimensions.



### Kernels

- The main shortcoming of PCA is that it is unable to capture nonlinear structures in the data.
- Consider the following example of linearly inseparable 1-D set of examples and then extracting polynomial (second order) features:


```
In [2]: def plot_kernel_example():
x_1 = np.random.randint(0,6, size=(3,))
x_2 = np.random.randint(8,14, size=(3,))
x_3 = np.random.randint(20,25, size=(3,))

x_1_p = x_1 ** 2
x_2_p = x_2 ** 2
x_3_p = x_3 ** 2

x_class = np.linspace(0, 30, 400)
y_class = 25 * x_class - 130

fig = plt.figure(figsize=(8,5))
ax_1 = fig.add_subplot(211)
ax_1.scatter(x_1, np.zeros_like(x_1), marker='*', color='b')
ax_1.scatter(x_2, np.zeros_like(x_2), marker='^', color='r')
ax_1.scatter(x_3, np.zeros_like(x_3), marker='*', color='b')
ax_1.grid()
ax_2 = fig.add_subplot(212)
ax_2.scatter(x_1, x_1_p, marker='*', color='b')
ax_2.scatter(x_2, x_2_p, marker='^', color='r')
ax_2.scatter(x_3, x_3_p, marker='*', color='b')
ax_2.plot(x_class, y_class)
ax_2.grid()
```

```
In [3]: plot_kernel_example()
```



- Adding polynomial features is simple to implement and can work great with all sorts of ML algorithms.
- At a **low polynomial** degree it cannot deal with more complex datasets.
- At a **high polynomial** degree there are a lot of features, which makes the computation very slow.
  - Computation in the feature space can be costly because it is high dimensional (even go to infinity).
- **The Kernel Trick** comes to the rescue!
  - It makes it possible to get the same result as if you added many features (even in high dimension), **without actually adding them!**
    - So there is no computational disaster resulting from the large number of features.



## The Kernel Trick

- When the **data points**,  $\{x^{(i)}\}_{i=1}^m$ , only appear as **inner (dot) products** -  $(x^{(i)})^T x^{(j)}$ , we can apply the kernel trick.
- As long as we can calculate the inner product in the **feature space**, we don't need the mapping explicitly.
  - The kernel function, denoted  $K(x^{(i)}, x^{(j)}) = \phi(x^{(i)})^T \phi(x^{(j)})$ , stands for the inner product in the feature space, and can easily be plugged into the optimization.
  - If we could find a function  $K(x^{(i)}, x^{(j)})$  that is simple enough, we can actually save us the huge number of calculations required to calculate  $\phi(x^{(i)})^T \phi(x^{(j)})$ , or even prevent us from the feature extraction itself,  $\phi(x)$ , this is the **kernel trick**.

- There are no restrictions on the form of  $x^{(i)}, x^{(j)}$ , but there are restrictions on the kernel function  $K(x^{(i)}, x^{(j)})$ 
  - $K(x^{(i)}, x^{(j)})$  must satisfy the **Mercer Condition** - the  $n \times n$  kernel matrix  $K(x^{(i)}, x^{(j)})$  is **PSD**
  - This implies that the quadratic optimization is convex
    - It has a unique solution and can be solved efficiently
- Mercer's Theorem:** Suppose that  $k(x, y)$  is a symmetric real valued function such that  $k(x, y) \leq C$  for some  $C > 0$ . Suppose the linear operator

$$\mathcal{L}(f)(x) = \int_{\mathcal{R}^D} k(x, y) f(y) dy$$

is PSD, that is

$$\int \int_{\mathcal{R}^D} k(x, y) f(y) dy dx \geq 0$$

Let  $\psi_i$  be the normalized orthogonal eigenfunctions of  $\mathcal{L}$ , associated with eigenvalues  $\lambda_i > 0$ , sorted in non-decreasing order, and let  $m$  be the number of non-zero eigenvalues. Then, the sequence of eigenvalues is summable, namely  $\sum_i |\lambda_i| < \infty$ , and the kernel can be expressed as

$$k(x, y) = \sum_{i=1}^m \lambda_i \psi_i(x) \psi_i(y)$$

We thus see that by choosing  $\phi_i(x) = \sqrt{\lambda_i} \psi_i(x)$ , a feature based representation is obtained.



## The Kernel Trick - Polynomial Kernel Example

- We will now demonstrate the kernel trick for a  $2^{nd}$ -degree polynomial mapping
- Suppose the feature mapping (or extraction) is defined as follows:

$$\phi(u) = \phi\left(\begin{bmatrix} u_1 \\ u_2 \end{bmatrix}\right) = \begin{bmatrix} 1 \\ \sqrt{2}u_1 \\ \sqrt{2}u_2 \\ \sqrt{2}u_1u_2 \\ u_1^2 \\ u_2^2 \end{bmatrix}$$

- The inner product  $\phi(u)^T \phi(v)$ :
$$\phi(u)^T \phi(v) = 1 + 2u_1v_1 + 2u_2v_2 + 2u_1u_2v_1v_2 + u_1^2v_1^2 + u_2^2v_2^2 = (1 + u_1v_1 + u_2v_2)^2 = (1 + u^T v)^2$$
  - HOW ABOUT THAT?** - We just found out that we can calculate  $\phi(u)^T \phi(v)$  just by calculating the input inner product  $u^T v$  and plugging it in  $(1 + u^T v)^2$ , thus saving us the whole feature mapping calculation process!
- We define the  $2^{nd}$  **polynomial kernel**:
$$K(u, v) = (1 + u^T v)^2$$
  - No need to calculate  $\phi(u)^T \phi(v)$  explicitly!
  - The use of kernel function to avoid computing  $\phi(\cdot)$  explicitly is known as **the kernel trick**



## Other Kernel Functions

- Linear kernel:
- Polynomial kernel with degree  $d$ :

$$K(u, v) = u^T v$$

$$K(u, v) = (\alpha u^T v + \beta)^d$$

- The feature space is all monomials up to degree  $d$

- **Radial Basis Function** (Gaussian RBF) kernel with width  $\sigma$ :

$$K(u, v) = e^{-\frac{\|u-v\|^2}{2\sigma^2}} = e^{-\gamma\|u-v\|^2} = \exp(-\gamma\|u-v\|^2)$$

- It is also a type of a *similarity function* that measures *how far are two points?*
- RBF covers the space with "balls" of a fixed radius, where the centers are the support vectors (in SVM).
- The radius of the "balls" is determined by the parameter  $\gamma = \frac{1}{2\sigma^2}$ 
  - A **smaller**  $\gamma$  means a larger radius, a lower "model complexity"
  - A **larger**  $\gamma$  means a smaller radius, a finer grain coverage which may possibly lead to overfitting.
- The feature space is **infinite-dimensional** (hint for HW: expand the exponential to an infinite series using Taylor).
- Sigmoid kernel, with parameters  $\kappa$  and  $\theta$ :

$$K(u, v) = \tanh(\kappa u^T v + \theta)$$

- It does not satisfy the *Mercer* condition on all  $\kappa, \theta$

- Techniques for constructing new kernels (from *Pattern Recognition and Machine Learning*, C.Bishop):

Given valid kernels  $k_1(\mathbf{x}, \mathbf{x}')$  and  $k_2(\mathbf{x}, \mathbf{x}')$ , the following new kernels will also be valid:

$$k(\mathbf{x}, \mathbf{x}') = ck_1(\mathbf{x}, \mathbf{x}') \quad (6.13)$$

$$k(\mathbf{x}, \mathbf{x}') = f(\mathbf{x})k_1(\mathbf{x}, \mathbf{x}')f(\mathbf{x}') \quad (6.14)$$

$$k(\mathbf{x}, \mathbf{x}') = q(k_1(\mathbf{x}, \mathbf{x}')) \quad (6.15)$$

$$k(\mathbf{x}, \mathbf{x}') = \exp(k_1(\mathbf{x}, \mathbf{x}')) \quad (6.16)$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}') \quad (6.17)$$

$$k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}') \quad (6.18)$$

$$k(\mathbf{x}, \mathbf{x}') = k_3(\phi(\mathbf{x}), \phi(\mathbf{x}')) \quad (6.19)$$

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T \mathbf{A} \mathbf{x}' \quad (6.20)$$

$$k(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}_a, \mathbf{x}'_a) + k_b(\mathbf{x}_b, \mathbf{x}'_b) \quad (6.21)$$

$$k(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}_a, \mathbf{x}'_a)k_b(\mathbf{x}_b, \mathbf{x}'_b) \quad (6.22)$$

where  $c > 0$  is a constant,  $f(\cdot)$  is any function,  $q(\cdot)$  is a polynomial with nonnegative coefficients,  $\phi(\mathbf{x})$  is a function from  $\mathbf{x}$  to  $\mathbb{R}^M$ ,  $k_3(\cdot, \cdot)$  is a valid kernel in  $\mathbb{R}^M$ ,  $\mathbf{A}$  is a symmetric positive semidefinite matrix,  $\mathbf{x}_a$  and  $\mathbf{x}_b$  are variables (not necessarily disjoint) with  $\mathbf{x} = (\mathbf{x}_a, \mathbf{x}_b)$ , and  $k_a$  and  $k_b$  are valid kernel functions over their respective spaces.



### Exercise - Valid Kernels

Let  $k_1(x, y)$  and  $k_2(x, y)$  be valid kernels according to Mercer's theorem (that is, they represent an inner product in some transformed space). Prove that the following functions are valid kernels as well:

1. Prove:

$$k(x, y) = k_1(x, y) + k_2(x, y)$$

2. Prove:

$$k(x, y) = f(x)k_1(x, y)f(y) \text{ where } f: \mathcal{R}^D \rightarrow \mathcal{R}$$



### Solution

The basic trick here is using the initial knowledge that both kernels are valid according to Mercer's theorem.

### Section 1

- Prove for:

$$k(x, y) = k_1(x, y) + k_2(x, y)$$

Using Mercer's theorem we know that there exists mappings  $\phi_1(x)$  and  $\phi_2(x)$  such that:

$$k_1(x, y) = \phi_1(x)^T \phi_1(y)$$

$$k_2(x, y) = \phi_2(x)^T \phi_2(y)$$

Thus,

$$\begin{aligned} k(x, y) &= k_1(x, y) + k_2(x, y) \\ &= \phi_1(x)^T \phi_1(y) + \phi_2(x)^T \phi_2(y) \\ &= [\phi_1(x), \phi_2(x)]^T [\phi_1(y), \phi_2(y)] = \Phi(x)^T \Phi(y) \end{aligned}$$

- $[\phi_1(x), \phi_2(x)]$  is a concatenation of the vectors  $\phi_1(x), \phi_2(x)$ :

$$[\phi_{11}(x) \quad \dots \quad \phi_{1n}(x) \quad \phi_{21}(x) \quad \dots \quad \phi_{2m}(x)]^T$$

- Thus,  $k(x, y)$  represents an inner product in some transformed space.

## Section 2

- Prove for:

$$k(x, y) = f(x)k_1(x, y)f(y) \text{ where } f: \mathcal{R}^D \rightarrow \mathcal{R}$$

Again, using Mercer's theorem:

$$\begin{aligned} k(x, y) &= f(x)k_1(x, y)f(y) \\ &= f(x)\phi_1(x)^T\phi_1(y)f(y) \\ &= \Phi(x)^T\Phi(y) \end{aligned}$$

where  $\Phi(x) = f(x)\phi_1(x)$ .

- Recall that  $\phi_1(x)$  is a vector and  $f(x)$  is a **scalar**.
- Thus,  $k(x, y)$  represents an inner product in some transformed space.



## Nonlinear PCA - NLPCA (& Kernel PCA)

- The main shortcoming of PCA is that it is **unable to capture nonlinear structures in the data**.
- The basic idea behind Kernel PCA is that by embedding the data into a high-dimensional space for which linear sub-spaces capture the relevant structures.
- Mathematically:

$$x \in \mathcal{R}^D \xrightarrow{\text{nonlinear}} \phi(x) \in \mathcal{R}^m \xrightarrow{\text{linear}} y \in \mathcal{R}^d, m \geq D > d$$

(and usually  $m \gg D$ )

- The basic steps of Nonlinear PCA:
  - Using the mapping  $\phi(x)$ , we map each data point  $x_i \rightarrow \phi(x_i)$ ,  $i = 1, \dots, n$ .
  - As in regular PCA, we center (or normalize/standardize) the data. We define:

$$\begin{aligned} \Phi &= (\phi(x_1) - \bar{\phi}, \dots, \phi(x_n) - \bar{\phi}) \in \mathcal{R}^{m \times n} \\ \bar{\phi} &= \frac{1}{n} \sum_{j=1}^n \phi(x_j) \end{aligned}$$

- Define the sample covariance matrix:

$$\hat{P}_{\phi(x)} = \frac{1}{n} \sum_{j=1}^n (\phi(x_j) - \bar{\phi})(\phi(x_j) - \bar{\phi})^T = \frac{1}{n} \Phi \Phi^T \in \mathcal{R}^{m \times m}$$

- Perform the *eigen-decomposition* of the sample covariance matrix:

$$\hat{P}_{\phi(x)} u_i = \lambda_i u_i, i = 1, \dots, m$$

- The  $d$  *nonlinear* Principal Components (PCs):

$$y_i = u_i^T (\phi(x) - \bar{\phi}) \in \mathcal{R}, u_i \in \mathcal{R}^m, i = 1, \dots, d$$

- Computational cost:
  - Computing the eigenvectors of a  $m \times m$  matrix requires (in general)  $O(m^3)$  steps - very expensive for large  $m$ !
- **Note** - the features  $\phi(x_i)$  need to be centered even if the  $x_i$  are centered, since the mapping  $\phi$  is **nonlinear**.
- Reducing the computational cost using the **Transpose Trick**:
  - The dimension  $m$  of the feature space may be very large which may be very expensive computationally.
  - Using the *transpose trick* which is based on the relation between the eigen-decomposition of  $\Phi \Phi^T$  and  $\Phi^T \Phi$ , we can reduce the cost.
    - Recall the **eigenvalues** of  $\Phi \Phi^T$  and  $\Phi^T \Phi$  are equal.
  - Doing the same derivation as before (look up...) we get that if  $w$  is an eigen-vector of  $\Phi^T \Phi$  and  $v$  is an eigen-vector of  $\Phi \Phi^T$  then:
 
$$u = \lambda^{-1} \Phi w$$
    - It is much cheaper to compute as the number of samples  $n \ll m$ .
- The  $d$  *nonlinear* Principal Components (PCs):

$$y_i = u_i^T (\phi(x) - \bar{\phi}) = \lambda_i^{-1} w_i^T \Phi^T (\phi(x) - \bar{\phi}) \in \mathcal{R}, u_i \in \mathcal{R}^m, i = 1, \dots, d$$



## Kernel PCA

- Notice that the nonlinear PCs only require the computation of *inner products* between the features  $\phi$ . This indicates that we can use **the kernel trick** to calculate these inner products.
- We define the kernel and *centered kernel functions*:

$$k(x, y) = \phi(x)^T \phi(y)$$

$$\tilde{k}(x, y) = (\phi(x) - \bar{\phi})^T (\phi(y) - \bar{\phi})$$

- Note that:

$$\tilde{k}(x, y) = k(x, y) - \frac{1}{n} \sum_{j=1}^n k(x, x_j) - \frac{1}{n} \sum_{i=1}^n k(x_i, y) + \frac{1}{n^2} \sum_{i=1}^n \sum_{j=1}^n k(x_i, y_j)$$

- Define the kernel and centered *kernel matrices*, which are defined w.r.t. the data points:

$$K = [k_{ij}], \quad k_{ij} = k(x_i, x_j)$$

$$\tilde{K} = \Phi^T \Phi = [\tilde{k}_{ij}], \quad \tilde{k}_{ij} = \tilde{k}(x_i, x_j)$$

### Claim

The nonlinear PCs can be computed directly in terms of the matrix  $\tilde{K}$ .

- It can be shown that:

$$\tilde{K} = \left(I - \frac{1}{n} \mathbf{1}\mathbf{1}^T\right) K \left(I - \frac{1}{n} \mathbf{1}\mathbf{1}^T\right) = J K J$$

- $\mathbf{1} = (1, 1, \dots, 1)^T \in \mathcal{R}^n$

- We also define the vectors:

$$\mathbf{k}_x = [k(x_1, x), \dots, k(x_n, x)]^T$$

$$\tilde{\mathbf{k}}_x = \Phi^T (\phi(x) - \bar{\phi}) = [\tilde{k}(x_1, x), \dots, \tilde{k}(x_n, x)]^T$$

- The nonlinear PCs can then be computed as:

$$y_i = w_i^T \Phi^T (\phi(x) - \bar{\phi}) = w_i^T \tilde{\mathbf{k}}_x, \quad i = 1, 2, \dots, d$$

- $w_i$  are the **eigenvectors** of  $\tilde{K}$ :

$$\tilde{K} w_i = \lambda_i w_i$$

$$\|w_i\|^2 = \lambda_i^{-1}, \quad [w_1, \dots, w_d] = V_d \Lambda_d^{-\frac{1}{2}}$$

- $V_d$  and  $\Lambda_d$  are obtained from the top- $d$  eigenvectors of  $\tilde{K}$ 's eigen-decomposition:

$$\tilde{K} = V_{\tilde{K}} \Lambda_{\tilde{K}} V_{\tilde{K}}^T$$

- Concatenating the components  $y_i$  we get (the projection):

$$Y = (y_1, y_2, \dots, y_d)^T = \Lambda_d^{-\frac{1}{2}} V_d^T \tilde{K}$$

- Notice that  $\Lambda_d^{-\frac{1}{2}}$  is the normalization factor (recall that  $\Lambda_d$  is a diagonal matrix with the eigenvalues on the diagonal).

- The low-dimensional nonlinear representation is obtained from the top- $d$  eigenvectors of the matrix  $\tilde{K}$ .

- The problem has been solved in the high-dimensional feature space  $\phi$  while not explicitly computing  $\phi$  or even working in  $\mathcal{R}^m$

### NLPCA Algorithm (Kernel PCA - KPCA)

- **Input:** A set of points  $\mathcal{X} = \{x_1, x_2, \dots, x_N\} \subset \mathcal{R}^D$ , and a map  $\phi : \mathcal{R}^D \rightarrow \mathcal{R}^M$  or a symmetric positive definite kernel function  $k : \mathcal{R}^D \times \mathcal{R}^D \rightarrow \mathcal{R}$

1. Compute  $\bar{\phi} = \frac{1}{N} \sum_j \phi(x_j)$  and the centered embedded data matrix  $\Phi$  or the kernel matrix  $K$ .

2. Compute the centered kernel matrix:

$$\tilde{K} = \Phi^T \Phi \text{ or } \tilde{K} = \left(I - \frac{1}{n} \mathbf{1}\mathbf{1}^T\right) K \left(I - \frac{1}{n} \mathbf{1}\mathbf{1}^T\right) = J K J \text{ or } (\tilde{k}(x_i, x_j)) \in \mathcal{R}^{N \times N}$$

3. Compute the eigenvectors  $w_i \in \mathcal{R}^N$  of  $\tilde{K}$ :

$$\tilde{K} w_i = \lambda_i w_i$$

and normalize so that  $\|w_i\|^2 = \lambda_i^{-1}$  (can be done as in step 4).

- For every data point  $x$ , its  $i^{th}$  nonlinear principal component is given by:

$$y_i = w_i^T \Phi^T (\phi(x) - \bar{\phi}) \text{ or } w_i^T [\tilde{k}(x_1, x), \dots, \tilde{k}(x_N, x)]^T$$

for  $i = 1, 2, \dots, d$

4. The projection is given by:

$$Y = (y_1, y_2, \dots, y_d)^T = \Lambda_d^{-\frac{1}{2}} V_d^T \tilde{K}$$

- Output:** A set of points  $\{y_j\}_{j=1}^N$  lying in  $\mathcal{R}^d$ , where  $y_{ij}$  in the  $i^{th}$  nonlinear principal component of  $x_j$  for  $i = 1, \dots, d$  and  $j = 1, \dots, N$

```
In [4]: def plot_circles_example():
# example - circles
X, y = make_circles(n_samples=1000, random_state=123, noise=0.1, factor=0.2)
pca = PCA(n_components=2)
kpca = KernelPCA(n_components=2, kernel='rbf', gamma=15) # Radial Basis Function kernel
X_pca = pca.fit_transform(X)
X_kpca = kpca.fit_transform(X)

# plot
fig = plt.figure(figsize=(10,4))

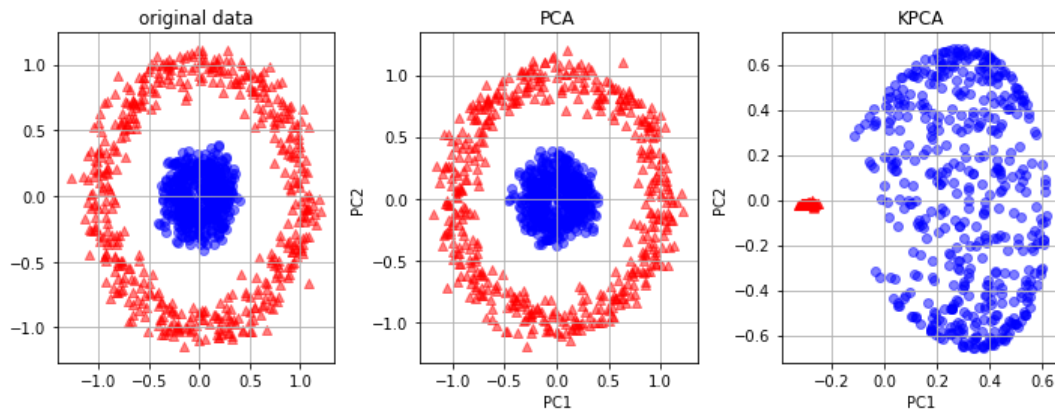
# original data
ax1 = fig.add_subplot(1,3,1)
ax1.scatter(X[y==0, 0], X[y==0, 1], color='red', marker='^', alpha=0.5)
ax1.scatter(X[y==1, 0], X[y==1, 1], color='blue', marker='o', alpha=0.5)
ax1.grid()
ax1.set_title("original data")

# pca
ax2 = fig.add_subplot(1,3,2)
ax2.scatter(X_pca[y==0, 0], X_pca[y==0, 1], color='red', marker='^', alpha=0.5)
ax2.scatter(X_pca[y==1, 0], X_pca[y==1, 1], color='blue', marker='o', alpha=0.5)
ax2.grid()
ax2.set_title("PCA")
ax2.set_xlabel("PC1")
ax2.set_ylabel("PC2")

# kpca
ax3 = fig.add_subplot(1,3,3)
ax3.scatter(X_kpca[y==0, 0], X_kpca[y==0, 1], color='red', marker='^', alpha=0.5)
ax3.scatter(X_kpca[y==1, 0], X_kpca[y==1, 1], color='blue', marker='o', alpha=0.5)
ax3.grid()
ax3.set_title("KPCA")
ax3.set_xlabel("PC1")
ax3.set_ylabel("PC2")

plt.tight_layout()
```

In [5]: `plot_circles_example()`



```
In [6]: def plot_moons_example():
# example - moons
X, y = make_moons(n_samples=100, random_state=123)
pca = PCA(n_components=2)
kpca = KernelPCA(n_components=2, kernel='rbf', gamma=15) # Radial Basis Function kernel
X_pca = pca.fit_transform(X)
X_kpca = kpca.fit_transform(X)

# plot
fig = plt.figure(figsize=(10,4))


# original data
ax1 = fig.add_subplot(1,3,1)
ax1.scatter(X[y==0, 0], X[y==0, 1], color='red', marker='^', alpha=0.5)
ax1.scatter(X[y==1, 0], X[y==1, 1], color='blue', marker='o', alpha=0.5)
ax1.grid()
ax1.set_title("original data")

# pca
ax2 = fig.add_subplot(1,3,2)
ax2.scatter(X_pca[y==0, 0], X_pca[y==0, 1], color='red', marker='^', alpha=0.5)
ax2.scatter(X_pca[y==1, 0], X_pca[y==1, 1], color='blue', marker='o', alpha=0.5)
ax2.grid()
ax2.set_title("PCA")
ax2.set_xlabel("PC1")
ax2.set_ylabel("PC2")

# kpca
ax3 = fig.add_subplot(1,3,3)
ax3.scatter(X_kpca[y==0, 0], X_kpca[y==0, 1], color='red', marker='^', alpha=0.5)
ax3.scatter(X_kpca[y==1, 0], X_kpca[y==1, 1], color='blue', marker='o', alpha=0.5)
ax3.grid()
ax3.set_title("KPCA")
ax3.set_xlabel("PC1")
ax3.set_ylabel("PC2")

plt.tight_layout()
```

In [7]: `plot_moons_example()`







## Recommended Videos



### Warning!

- These videos do not replace the lectures and tutorials.
- Please use these to get a better understanding of the material, and not as an alternative to the written material.

### Video By Subject

- The Kernel Trick - [Machine Learning Fundamentals - The Kernel Trick \(https://www.youtube.com/watch?v=JiM\\_LXpAtLc\)](https://www.youtube.com/watch?v=JiM_LXpAtLc).
- Kernel PCA - [David Thompson: Nonlinear Dimensionality Reduction: KPCA \(https://www.youtube.com/watch?v=HbDHohXPLnU\)](https://www.youtube.com/watch?v=HbDHohXPLnU).



## Credits

- Icons from [Icons8.com \(https://icons8.com/\)](https://icons8.com/) - [https://icons8.com \(https://icons8.com\)](https://icons8.com/)
- Datasets from [Kaggle \(https://www.kaggle.com/\)](https://www.kaggle.com/) - [https://www.kaggle.com/ \(https://www.kaggle.com/\)](https://www.kaggle.com/)